

EECS 496: Sequential Decision Making

Soumya Ray

sray@case.edu

Office: Olin 516

Office hours : T 4-5:30 or by appointment

Recap

- We can handle evidence in approximate inference in several ways. The simplest is to do _____ sampling, but _____ samples which _____ with the evidence. This is called _____.
- Alternatively, we could just sample the _____ variables. Then assign each sample a _____, which is the _____ of the _____ given _____. This is called _____.
- If we have evidence _____ in the topological sort, or evidence that is _____, both the above have problems. These problems are (i) _____ (ii) _____.
- To alleviate these issues we can stop _____ samples. This creates a _____ over samples. The resulting algorithm is called _____.
- What is burn in time? Why do we need it?
- A specific approach in the case of Bayes nets is called _____. This generates the next sample by choosing a _____ variable V . Then it samples from $\Pr(V|\text{_____})$.
- If a Markov chain is _____ and satisfies _____, then it will eventually converge to a _____.
- Using this, we can show that Gibbs sampling produces the right result for a Bayes net, because _____.

Today

- Sequential Probabilistic Models (Ch 15)

Sequential Models

- So far, no explicit representation of time
- What happens when we have a random process evolving over time?

Sequential Models

- We see a sequence of observations o_1, \dots, o_n
- Each element is drawn from some background alphabet or **vocabulary**
 - Text classification---each observation is a word
 - PFM---each observation is an amino acid
 - Parsing---each observation is a word or phrase
 - Activity recognition from video---each observation is a frame
 - Radar/Lidar/Sensing---each observation is a sensor measurement
- Observation sequences have varying lengths

Generative Process Model

- Assume sequential data is generated by an underlying generating process
- This process has a **state** that could be **discrete or continuous**
- The state evolves over time
- At **discrete time points**, we observe something about the state
 - These are our sequences of observations

Discrete Time, Discrete State Sequential Probabilistic Models

Minus 1th approach

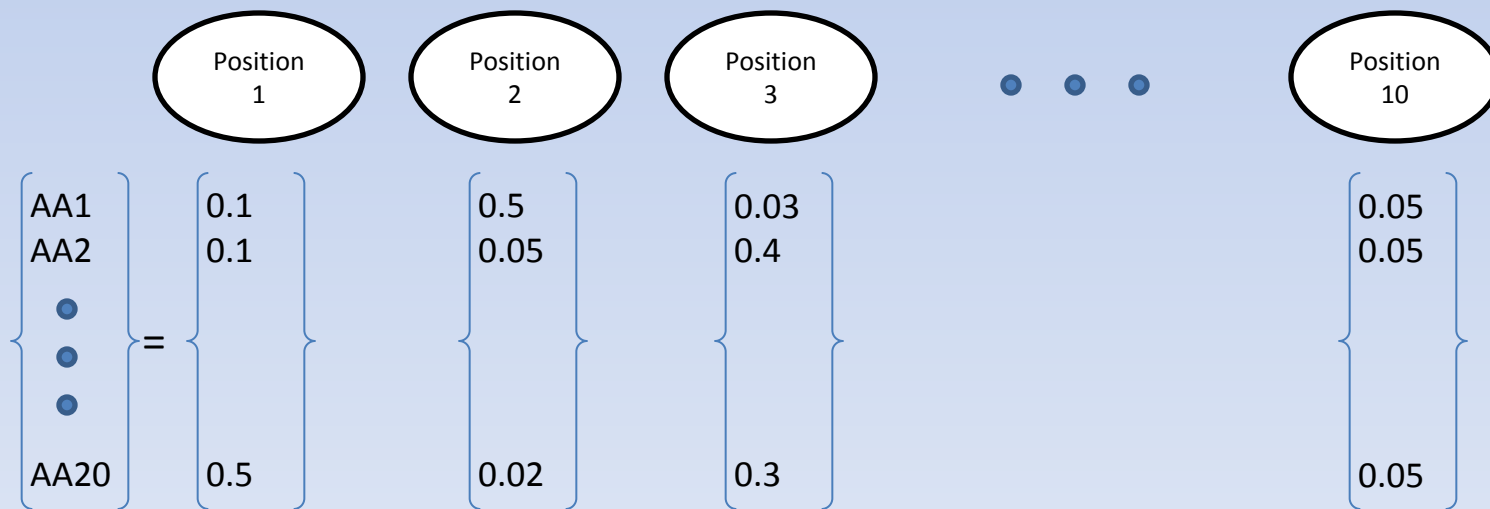
- Let's ignore everything about process state and dependence; pretend the data isn't sequential at all
 - Sometimes a reasonable first approximation
- Naïve Bayes for text classification

Zeroth approach

- Let's ignore dependence, but not process state
- We'll record a probability distribution over observations at each state of interest
 - “Position Specific Scoring Matrix”
- Only useful for fixed length sequences
 - Can also be used to find the subsequence of length k with highest probability

Example

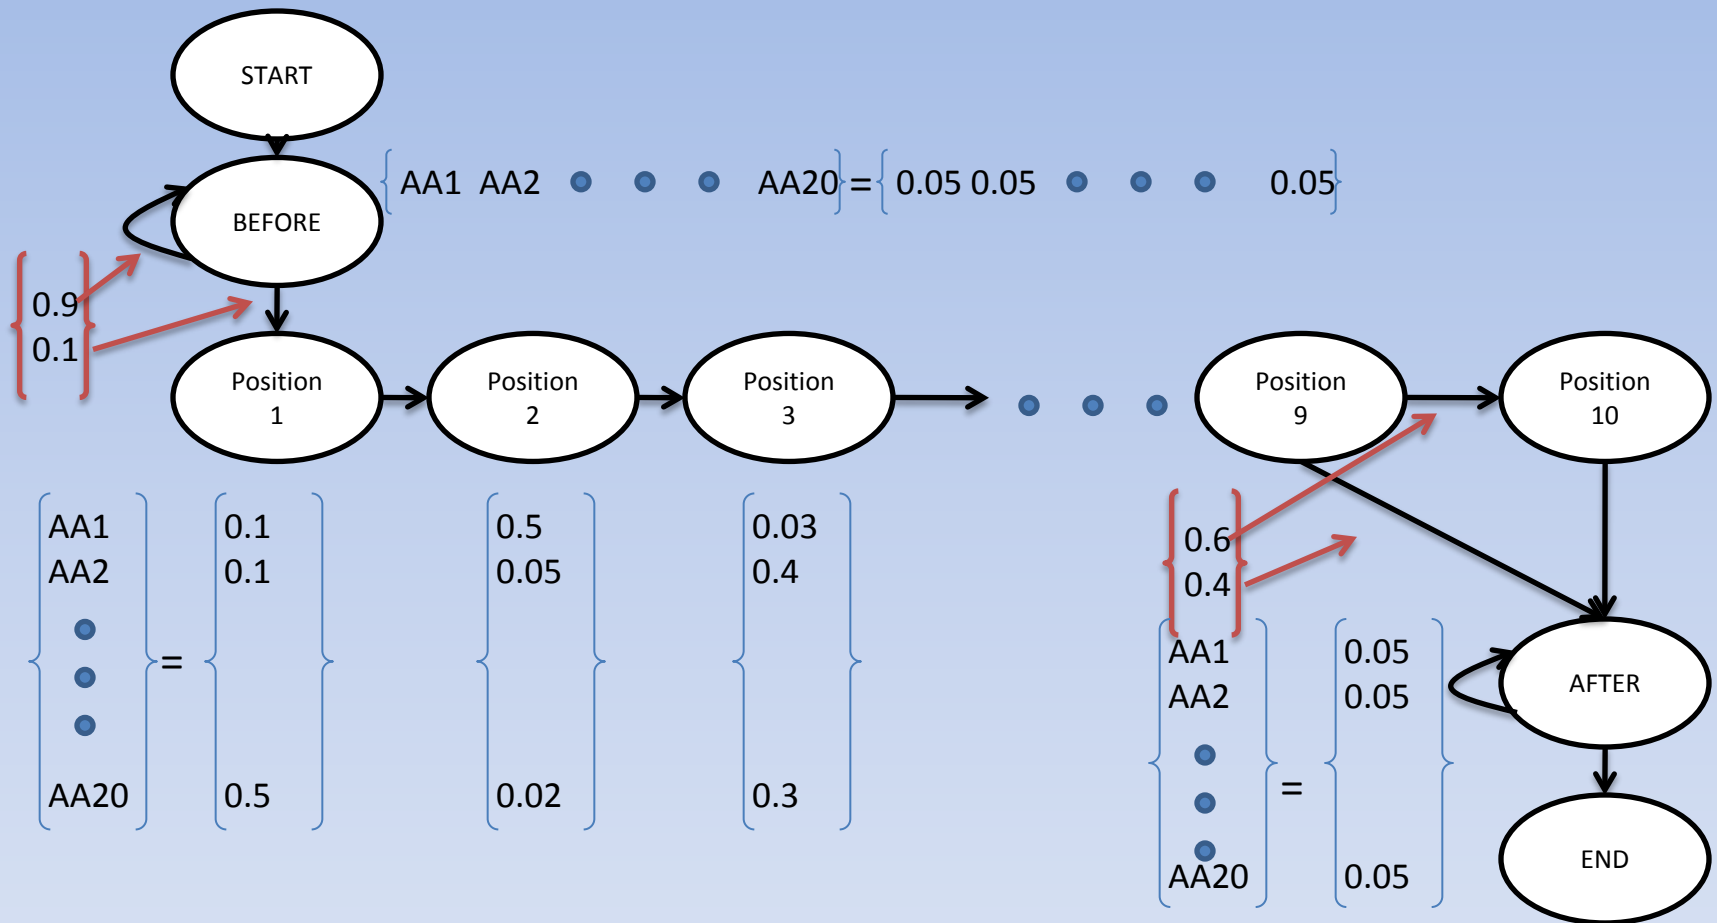
- Suppose members of some family of proteins have a *motif*: a sequence of 10 amino acids somewhere that is specific to this family



k^{th} order approach

- We'll model dependence of “order k ”
 - Assume each state is dependent on the previous k states
- We'll study the case for $k=1$
 - $k>1$ are straightforward extensions (just messy algebraically)
 - “First order Hidden Markov Models”

Hidden Markov Models



Set of states = {BEFORE, Position1, ..., Position10, AFTER}
Set of emissions = {AA1, AA2, ..., AA20}

Hidden Markov Models

- HMMs are generative process models for the **joint distribution $\Pr(\mathbf{s}, \mathbf{o})$**

$$\Pr(\{s_1, s_2, \dots, s_n\}, \{o_1, o_2, \dots, o_n\}) = \Pr(s_1) \Pr(o_1 | s_1) \prod_{r=2}^n \Pr(o_r | s_r) \Pr(s_r | s_{r-1})$$

“Emission Probability”:
How likely is process state s_r to emit observation o_r

“Transition Probability”:
How likely is process state s_{r-1} to transition to state s_r

Questions

- What's “Markov” about this model?
- What changes for $k > 1$?
- What's “hidden”?

Key Issues

- Inference
 - What is the probability of an observed sequence \mathbf{o} ?
 - What is the most likely sequence of process states \mathbf{s} that could have emitted an observed sequence \mathbf{o} ?
- Learning
 - Given a training set of observation sequences and a model structure, how do we estimate parameters for the model?

Issue #1: $\Pr(\mathbf{o})$

- Clearly, $\Pr(\mathbf{o}) = \sum_{\mathbf{s}} \Pr(\mathbf{s}, \mathbf{o})$
 - Sum over all possible state sequences that could generate \mathbf{o}
 - But the number of state sequences that could generate \mathbf{o} could be exponential in the length of \mathbf{o}
- **Key observation:**
 - Many state sequences share prefixes
 - We only need to compute probabilities for a shared prefix *once*
 - In fact, because of the Markov property, we can do better
 - We can use dynamic programming to store and reuse these computations

Forward Algorithm

- Let $\alpha_k(i) = \Pr(o_1, \dots, o_i, s_i = k)$
 - Denotes the probability that the model has emitted the first i observations and is now in state k
- We want to compute $\alpha_{END}(n)$ (recall the observation sequence is extended with dummy START and END symbols)
- Construct a table of size n -by- m , n =length of observed sequence, m =number of states
- The forward algorithm is a dynamic programming procedure that will fill in this table with α values

Forward Algorithm

- Initialize: $\alpha_{START}(0) = 1, \alpha_k(0) = 0, k \neq START$

- Recursion: Emitting observation i Transition to state k

$$\alpha_k(i) = \Pr(o_i \mid s_i = k) \sum_p \alpha_p(i-1) \Pr(s_i = k \mid s_{i-1} = p)$$

$$\alpha_k(i) = \Pr(o_1, \dots, o_i, s_i = k)$$

$$\alpha_p(i-1) = \Pr(o_1, \dots, o_{i-1}, s_{i-1} = p)$$

Issue #2: Most Likely Path

- Given an observation sequence, what is the most likely sequence of states that could emit it?

$$\mathbf{s}^* = \arg \max_{\mathbf{s}} \Pr(\mathbf{s} \mid \mathbf{o})$$

- Dumb way: enumerate all possible \mathbf{s}
- Smart way: dynamic programming, as before, taking advantage of Markov property
 - Viterbi algorithm

Viterbi Algorithm

- Let $\gamma_k(i) = \Pr(o_1, \dots, o_i, s_i^* = k)$
 - Denotes the probability that the most likely path is at state k after emitting the first i observations
- We want $\gamma_{END}(n)$
- Notice that this just gives us the *probabilities*
 - To get the path, we will also need to maintain pointers to certain table elements

Viterbi Algorithm

- Initialize: $\gamma_{START}(0) = 1, \gamma_k(0) = 0, k \neq START$

- Recursion: Emitting observation i Transition to state k

$$\gamma_k(i) = \Pr(o_i \mid s_i^* = k) \max_p \gamma_p(i-1) \Pr(s_i^* = k \mid s_{i-1}^* = p)$$

$$\gamma_k(i) = \Pr(o_1, \dots, o_i, s_i^* = k)$$

$$\gamma_p(i-1) = \Pr(o_1, \dots, o_{i-1}, s_{i-1}^* = p)$$

To get the path, store the arg max's of the recursive computation.