

# EECS 496: Sequential Decision Making

Soumya Ray

[sray@case.edu](mailto:sray@case.edu)

Office: Olin 516

Office hours : T 4-5:30 or by appointment

# Recap

- We estimate model parameters from data using the technique of \_\_\_\_\_. This finds parameters that \_\_\_\_\_.
- For HMMs, the parameters are \_\_\_\_\_.
- Suppose the observations are annotated with state sequences. Then the emission MLEs are? The transition MLEs?
- If the observations are NOT annotated with state sequences, we use a general technique, called \_\_\_\_\_.
- In this technique we first \_\_\_\_\_. Then we \_\_\_\_\_ two steps. In the \_\_\_\_\_ step, we find the \_\_\_\_\_ of the missing information (or sometimes the \_\_\_\_\_ the missing information). In the \_\_\_\_\_ step, we find the \_\_\_\_\_ given the previous. Eventually this procedure \_\_\_\_\_ to a \_\_\_\_\_.
- When applied to HMMs this is called the \_\_\_\_\_ algorithm.

# Today

- Sequential Probabilistic Models (Ch 15)

# Baum-Welch Algorithm

- Start by randomly initializing all the emission and transition probabilities
- Proceed in two steps until (local) convergence
  - **E step**: Since we don't know the state sequences, we calculate the *expected* number of times each transition or emission is used (as if annotating observations with *possible* state sequences and averaging over them)
  - **M step**: Update the emission and transition probabilities using the expected counts (same as in the fully observed case above)

# Baum-Welch Algorithm (E step)

Expected number of times  $b$  is emitted by state  $k$  according to our training sample:  $E_{k;b}$

For each training sequence, for each time step where  $b$  is present, find the probability  $k$  was used at that point; add up all the probabilities

$O^1 = o_1 \ o_2 \ o_3$   
 $O^2 = o_1 \ o_2 \ \mathbf{b}$   
 $O^3 = o_1 \ o_2 \ o_3$   
 $O^4 = \mathbf{b} \ o_2 \ o_3$   
 $O^5 = o_1 \ o_2 \ o_3$   
 $O^6 = \mathbf{b} \ \mathbf{b} \ \mathbf{b}$

**b**

Pr( $S=k$  emitted this symbol)?

$\dots$   $O_n$   
 $\dots$   $O_m$   
 $\dots$   $O_k$   
 $\dots$   $\mathbf{b}$   $O_p$   
 $\dots$   $O_n$   
 $\dots$   $O_q$

# Baum-Welch Algorithm (E step)

Expected number of times  $b$  is emitted by state  $k$  according to our training sample:  $E_{k;b}$

For each training sequence, for each time step where  $b$  is present, find the probability  $k$  was used at that point; add up all the probabilities

$$\begin{aligned} E_{k;b} &= \sum_{j \in \text{sequences}} \sum_{i | o_i^j = b} \Pr(s_i = k | \mathbf{o}^j) \\ &= \sum_{j \in \text{sequences}} \frac{1}{\Pr(\mathbf{o}^j)} \sum_{i | o_i^j = b} \Pr(s_i = k, \mathbf{o}^j) \end{aligned}$$

# The Backward Algorithm

- Let  $\beta_k(i) = \Pr(o_{i+1}, \dots, o_n \mid s_i = k)$
- Calculated similarly to  $\alpha_k(i)$

- Then notice that:

$$\Pr(s_i = k, \{o_1, \dots, o_n\})$$

Prob of emitting sequence AND *using state k at  $i^{\text{th}}$  time step*

$$= \Pr(\{o_1, \dots, o_i\}, s_i = k) \Pr(\{o_{i+1}, \dots, o_n\} \mid \{o_1, \dots, o_i\}, s_i = k)$$

$$= \Pr(\{o_1, \dots, o_i\}, s_i = k) \Pr(\{o_{i+1}, \dots, o_n\} \mid s_i = k)$$

$$= \alpha_k(i) \beta_k(i)$$

# Baum-Welch Algorithm (E step)

Expected number of times  $b$  is emitted by state  $k$  according to our training sample:

For each training sequence, for each time step where  $b$  is present, find the probability  $k$  was used at that point; add up all the probabilities

$$\begin{aligned} E_{k;b} &= \sum_{j \in \text{sequences}} \sum_{i | o_i^j = b} \Pr(s_i = k | \mathbf{o}^j) \\ &= \sum_{j \in \text{sequences}} \frac{1}{\Pr(\mathbf{o}^j)} \sum_{i | o_i^j = b} \Pr(s_i = k, \mathbf{o}^j) = \sum_{j \in \text{sequences}} \sum_{i | o_i^j = b} \frac{\alpha_{k,j}(i) \beta_{k,j}(i)}{\alpha_{\text{END},j}(n)} \end{aligned}$$



# Baum-Welch Algorithm (E step)

Expected number of times a transition  $k \rightarrow l$  is used:  
For each training sequence, for each time step, find the probability  $k \rightarrow l$  was used at that point; add up all the probabilities

$$\begin{aligned} & \Pr(s_i = k, s_{i+1} = l \mid \mathbf{o}) \\ &= \frac{\alpha_k(i) \Pr(s_{i+1} = l \mid s_i = k) \Pr(o_{i+1} \mid s_{i+1}) \beta_l(i+1)}{\Pr(\mathbf{o})} \end{aligned}$$

$$C_{kl} = \sum_{j \in \text{sequences}} \frac{1}{\Pr(\mathbf{o}^j)} \sum_i \alpha_{k,j}(i) \Pr(s_{i+1} = l \mid s_i = k) \Pr(o_{i+1} \mid s_{i+1}) \beta_{l,j}(i+1)$$

# Baum-Welch Algorithm (M-Step)

- Emission distribution

$$\Pr(a \mid S = k) = \frac{E_{k;a}}{\sum_b E_{k;b}}$$

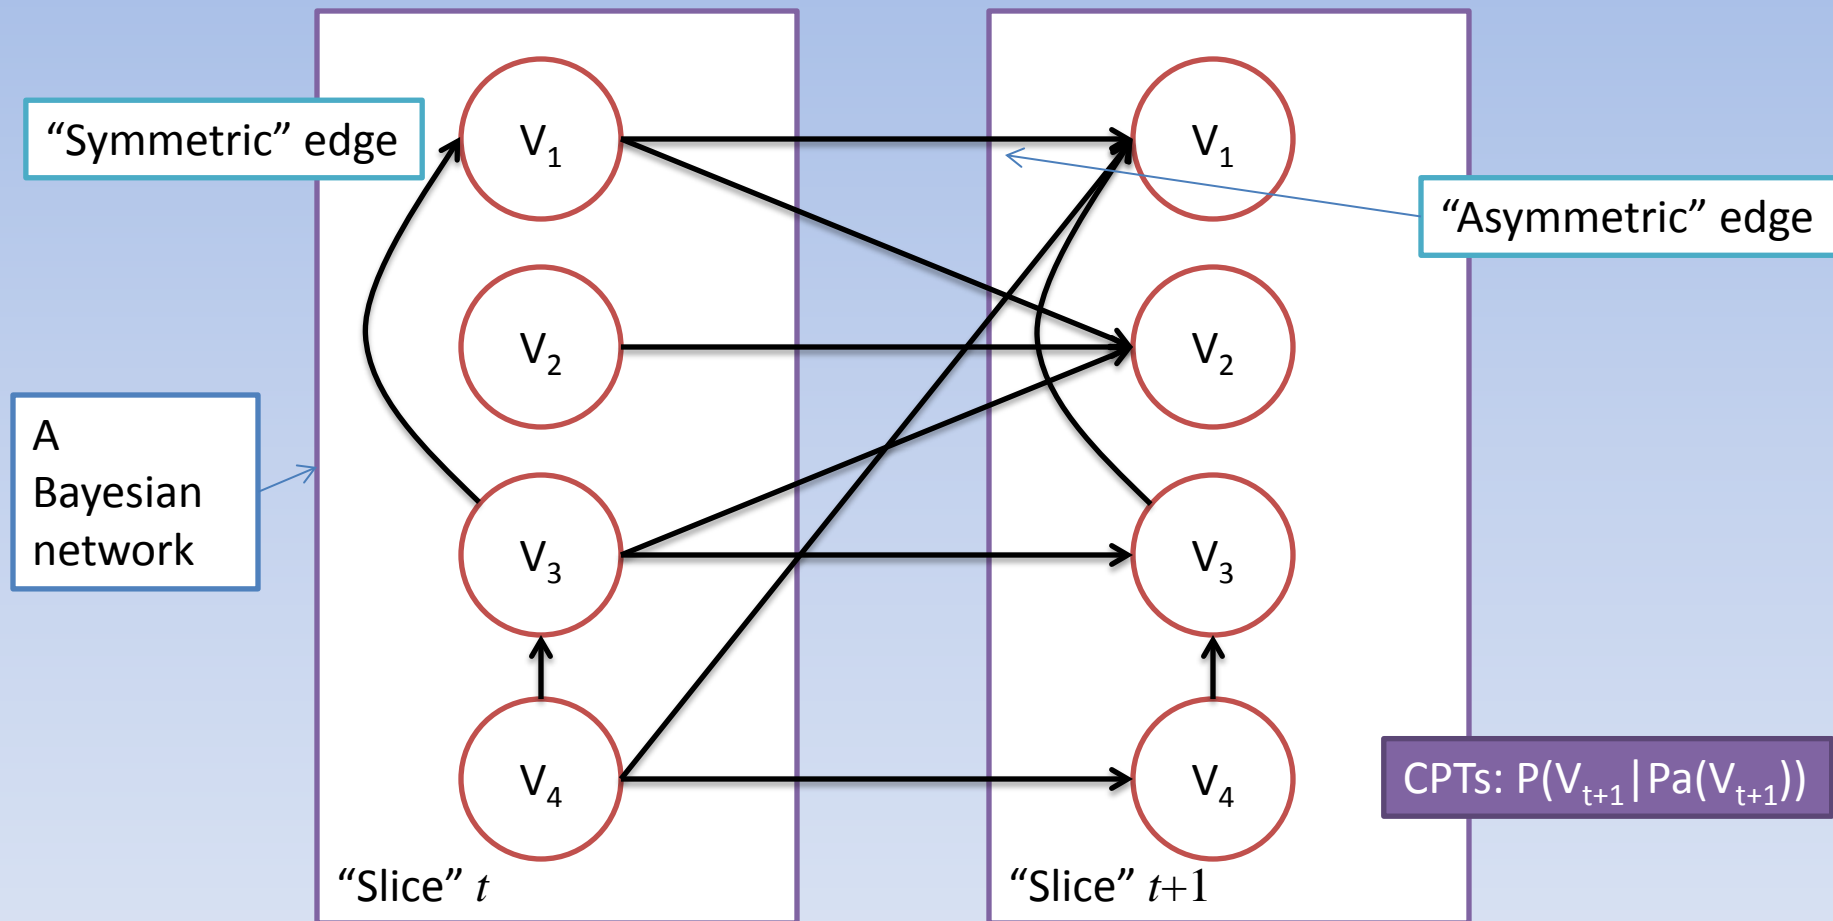
- Transition Distribution

$$\Pr(S = k \mid S = j) = \frac{C_{jk}}{\sum_l C_{jl}}$$

# Dynamic Bayesian Networks

- HMMs have *atomic* states
- In most problems, states have structure and the observations are some aspect of this structure
- Sequential probabilistic models with *factored* state representations are called Dynamic Bayesian Networks (DBNs)

# DBN Architecture



# Inference in DBNs

- Since a DBN is a BN, we could technically answer queries by “unrolling” it and applying variable elimination
- However this is generally computationally expensive and becomes increasingly infeasible over time
- Therefore we must use approximate inference algorithms

# Particle Filtering/ Sequential Monte Carlo

- An approach that extends likelihood weighting to sequential probabilistic models
- Problem: Must avoid “unrolling” probabilistic models
- So at each time step, need to represent the joint distribution over the variables, then propagate to the next step

# Particle Filtering/ Sequential Monte Carlo

- In PF, we adopt a nonparametric representation: *use a set of weighted samples* (“particles”) to approximate the underlying distribution
- Parameter: Sample size set  $N$

# Particle Filtering

1. Pick an initial set of  $N$  samples from a prior over the initial state distribution  $S_0$
2. Each time step:
  1. Propagate each sample  $i$  forward by sampling the next state from  $P(V_{t+1}^i | \text{Pa}(V_{t+1}^i))$ .
  2. Weight the sample by the likelihood it assigns to any evidence at  $t+1$ :  $w^i = P(e_{t+1} | \mathbf{V}_{t+1})$
  3. **Resample** a set of particles from the weighted sample in step 2. This set is unweighted.



# Particle Filtering

- It can be shown that the procedure is consistent (see book): the fraction of samples in any state is proportional to the true probabilities for being in that state (for large enough  $N$ )