

EECS 496: Sequential Decision Making

Soumya Ray

sray@case.edu

Office: Olin 516

Office hours : T 4-5:30 or by appointment

Recap

- With Kalman filters, the state is modeled as a set of ____ ____.
- The transition/sensor distributions for a new state is obtained as ____ functions of the old state with ____ ____.
- In continuous state scenarios, we are often interested in ____ prediction problems. These are where we ____.
- What is filtering? Smoothing? “Prediction”? Most likely path?
- In filtering, the state distribution at $(t+1)$ is a function of the __ __, the __ __ and the ____.
- If the sensor and transition distributions are __ __, then the ____ of filtering is ____.
- KFs make strong assumptions, which can be relaxed by integrating them with _____. In this case each _____ has a _____ distribution over its _____. Further flexibility could be added through _____ which have _____ distributions in the _____.

Today

- Part 2: Automated Planning (Ch 10 R&N)

Automated Planning

- Consider a situation where an agent has to carry out a sequence of actions to achieve a goal
- Suppose the agent starts off with detailed, *structured* knowledge of the world
 - Could we take advantage of this?

The Planning Problem

- Given:
 - An initial state of the world, described as a set of logical facts
 - A set of goal conditions, described as a set of logical facts
 - A set of actions, also described in logic
- Find a sequence of actions that will move the world from the initial state to the final state
 - This sequence is called a *plan*
 - Often also try to optimize some criteria


“Classical” Planning

- We’ll study planning algorithms designed to work when the world is:
 - Static
 - Deterministic
 - Fully observable
 - Discrete
 - Actions are instantaneous
- These restrictions can be relaxed (more or less)

Makespan

- Typically, in planning, we are interested in minimizing the *duration* of the plan
 - Equivalent to the number of actions in the current setup
 - This is called the *makespan*
- Nonclassical planning allows arbitrary plan metrics to be minimized

Example: Blocks World



Task: Starting with initial configuration of blocks, produce a desired goal configuration by moving blocks around.

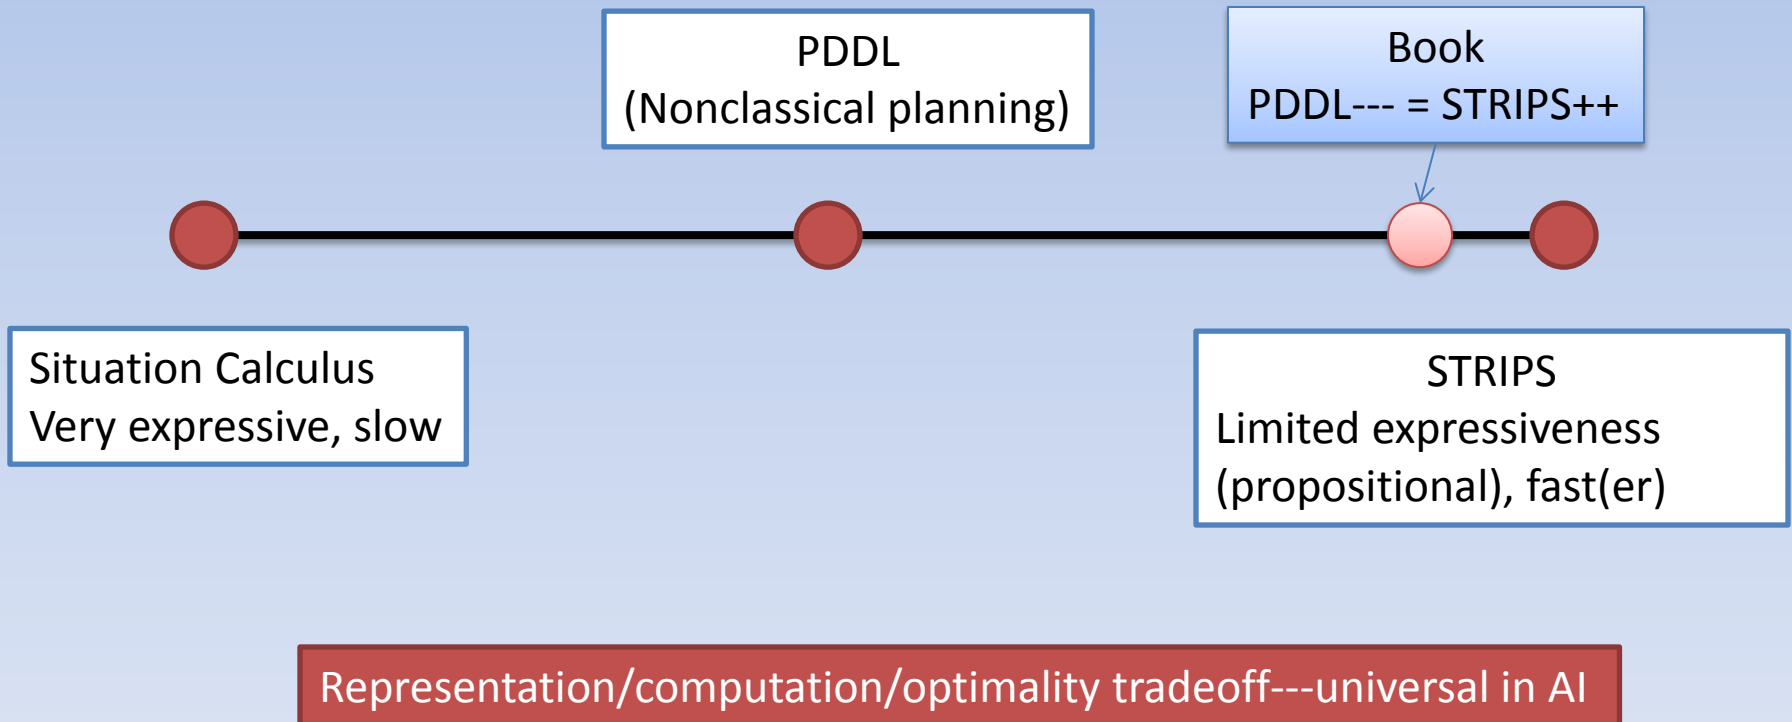
Extended Syntax of Planning

- Holding(A)
 - A “predicate” that can take multiple “arguments” and evaluates to true or false
 - Internally map to a proposition like Holding_A
- Holding(x)
 - This predicate takes a *variable* argument
 - When planning, the inference algorithm will try *all possible values for x* that are applicable (later)
 - Map to Holding_A, Holding_B,...

Extended Syntax of Planning

- The fragment of FOL that we just saw is called STRIPS (“Stanford Research Institute Problem Solver”)
- States, actions and goals will be represented in this language
 - Then we’ll see planning algorithms (which are inference algorithms in disguise) that find plans in this language

Extended Syntax of Planning



Representing States in STRIPS

- States in STRIPS are conjunctions of unnegated ground literals
 - All conditions that hold in that state
 - *Block(A), Block(B), On(A,B), On(B, Table), GripperEmpty*
 - Technical condition: also need to be “function free”
 - A function is an FOL construction that maps one object to another, e.g. $\text{Mother}(X)=Y$
 - The “Closed World Assumption” is used

No variables

Closed World Assumption (CWA)

- Anything that is not explicitly listed as part of a state definition is false
 - No “unknown” variables

Representing Goals in STRIPS

- Goals are conjunctions of unnegated, ground (function-free) literals
- Goals may not fully determine a state of the world
 - In this case, the goal is any state where these literals hold
- Example: $On(A,E) \wedge On(B,D)$

Representing Actions in STRIPS

- Want to represent an action of picking up a block from the table

Pickup_from_Table(x)

Preconditions: *Block(x), GripperEmpty, Clear(x), On(x, Table)*

Add List: *Holding(x)*

Delete List: *GripperEmpty, On(x, Table)*

“Applicability”: action can be used at a state iff its preconditions are satisfied

Representing Actions in STRIPS

- An “action schema” represents a non-ground action using three parts:
 - The action name and argument list
 - The **preconditions**: a conjunction of unnegated non-ground (function-free) literals. Any variables in this list are parameters to the action.
 - The **effects**: a conjunction of function-free literals describing how the state changes.



Can have variables

Add and Delete Lists

- Often, the unnegated literals in the action effects are collected into an “ADD” list, and the negated literals are collected into a “DELETE” list
 - Idea: Starting with initial state, to get result of applying action, add the literals in ADD list and delete the literals in the DELETE list


The STRIPS assumption

- Every possible effect of actions are listed
 - i.e., if a literal does not appear in the effects list, it is unchanged in the resulting state
 - Together with CWA, solves the “frame problem”
 - In general logic, a literal whose truth value is unspecified becomes “unknown”

Restrictions in STRIPS

- States are described by unnegated ground (function-free) literals
- CWA + STRIPS assumption
- Ground conjunctive goals
- Conjunctive effects of actions

Example: Blocks World



Example

- $\text{Init}(\text{On}(A, \text{Table}) \wedge \text{On}(B, \text{Table}) \wedge \text{On}(C, \text{Table}) \wedge \text{Block}(A) \wedge \text{Block}(B) \wedge \text{Block}(C) \wedge \text{Clear}(A) \wedge \text{Clear}(B) \wedge \text{Clear}(C) \wedge \text{GripperEmpty})$
- $\text{Goal}(\text{On}(A, B))$
- $\text{Action}(\text{MoveToTable}(b, x),$
 - $\text{Preconditions}(\text{On}(b, x) \wedge \text{Clear}(b) \wedge \text{Block}(b) \wedge \text{Block}(x) \wedge \text{GripperEmpty})$
 - $\text{AddEffects}(\text{On}(b, \text{Table}) \wedge \text{Clear}(x))$
 - $\text{DelEffects}(\text{On}(b, x))$

Planning Algorithms

- Given a STRIPS representation of a classical planning problem, how do we solve it?
 - Since the world is static, deterministic, fully observable, we could use search
 - Remember that in this case, the search algorithm is actually performing logical inference

Kinds of Search for Planning

- Search algorithms for classical planning fall into two categories
 - “State space planners”: States of the search problem are states of the world; search operators are actions of the world
 - “Plan space planners”: States of the search problem are partial plans; search operators are modifications to the current partial plan

Forward State-Space Search

- “Progression” planning
- Setup:
 - States=world states (in STRIPS)
 - Initial state=given
 - Operators=*applicable* actions (in STRIPS)
 - Goal test=given (in STRIPS)
 - Operator costs=unit (minimize number of actions)

Forward State-Space Search

- We could apply any search algorithm, e.g. A^*
- The key differences are:
 - Only applicable actions need to be explored at a state
 - Getting the next state is done through the STRIPS specification of states and actions
 - Heuristics are based on planning ideas

Search Heuristics

- From any state, want to estimate the number of actions to search termination admissibly
- Two possibilities:
 - Relax the planning problem
 - Consider subproblems

Relaxed Plans

- There are different ways to arrive at a less constrained planning problem
- One way is to remove all DELETE effects from STRIPS actions
 - This is admissible (why?)
 - To estimate this cost, need to run an internal planning loop; but this is usually very fast

Subproblems

- The goal is a conjunction of literals
- We can generate subproblems by just considering a single literal at a time
 - “Subgoal Independence” (admissible)
- Combine with the max function